
Master Course – High-Performance Computing

Parallel Programming with OpenMP
Part I

Olaf Schenk
Institute of Computational Science
USI Lugano, Switzerland
October 16, 2019

Outline

- Introduction into OpenMP
- Programming and Execution Model
 - **Parallel regions: team of threads**
 - **Syntax**
 - **Data environment (part 1)**
 - **Environment variables**
 - **Runtime library routines**
 - **Exercise 1: Parallel region / library calls / privat & shared variables**
- Worksharing directives
 - **Which thread executes which statement or operation?**
 - **Synchronization constructs, e.g., critical regions**
 - **Nesting and Binding**
 - **Exercise 2: Pi**
- Data environment and combined constructs
 - **Private and shared variables, Reduction clause**
 - **Combined parallel worksharing directives**
 - **Exercise 3: Pi with reduction clause and combined constructs**
 - **Exercise 4: Heat**
- Summary of OpenMP API
- OpenMP Pitfalls & Optimization Problems

Multiprocessor System with Shared Memory

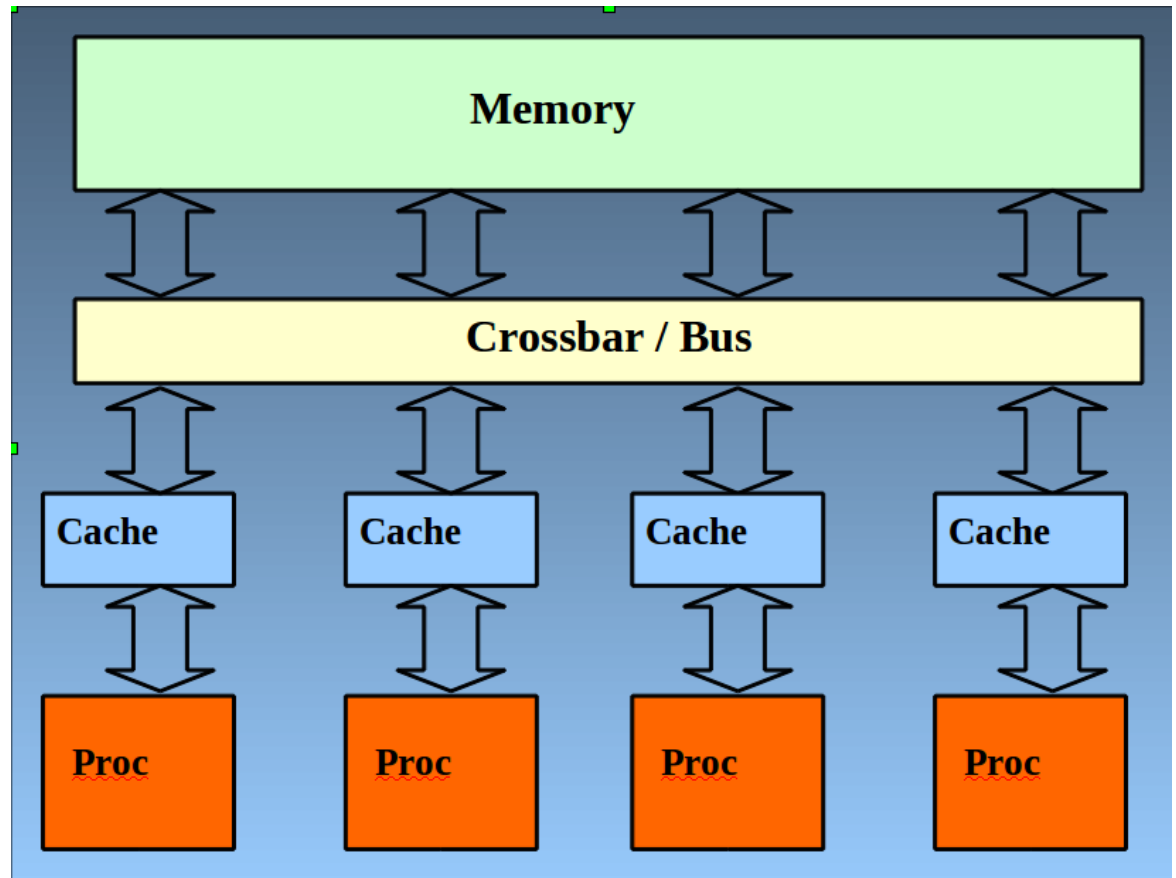
OpenMP

is a

**parallel programming
model**

for

**Shared-memory
multiprocessors**



OpenMP Overview: What is OpenMP?


- OpenMP is a standard programming model for shared memory parallel programming
- Portable across all shared-memory architectures
- It allows incremental parallelization
- Compiler based extensions to existing programming language
 - mainly by directives
 - a few library routines
- Fortran and C/C++ binding
- OpenMP is a standard

Parallel Computing: Effective Standards for Portable programming


- Thread Libraries
 - Win32 API
 - POSIX threads
- Compiler Directives
 - **OpenMP - portable shared memory parallelism**
 - OpenACC, OpenCL
- Global Address Space Languages
 - Unified Parallel C
 - Titanium
 - X10
 - Chapel
- Message Passing Libraries
 - MPI



Motivation: Why should I use OpenMP?



Where should I use OpenMP?



Multithreading versus Multi-Processing

- **Multiple Processes** (Heavyweight Process model)
 - traditional UNIX process model
 - interprocess communication techniques supported by OS: shared memory, sockets, file IO, memory map
 - Higher overhead associated with process creation and destruction
- **Multiple Threads** (Lightweight Process model, LWP)
 - thread concept: independent flow of control within one process with its own context: stack, register set
 - process data and opened files are shared
 - lower overhead of thread creation and destruction
 - shared address space
 - Auto-Parallelization, OpenMP, Explicit Multithreading using P-Threads
- **Hybrid Models** (e.g. MPI + OpenMP)


OpenMP - History

- 1997: OpenMP Version 1.0 for Fortran
 - de facto standard for shared memory programming
 - now available for all SMP systems
 - replaces proprietary parallelization directives and in many cases the explicit programming of [p]threads
 - 1998: OpenMP V1.0 for C and C++
 - 1999: OpenMP V1.1 for Fortran (error corrections, explanations)
 - 2000: OpenMP V2.0 for Fortran (support of Fortran90 modules)
 - 2001: OpenMP V2.0 for C and C++
 - 2008: OpenMP V3.0 for Fortran/ C and C++
 - 2014: OpenMP V4.0 for Fortran/C and C++



OpenMP - Information

- The OpenMP Architecture Review Board (ARB) Fortran and C Application Program Interfaces (APIs): www.openmp.org
- NCSA online course on OpenMP:
<http://www.citutor.org/login.php?course=24>
(you need to register for a free account, very useful)
- OpenMP tutorial from Lawrence Livermore National Laboratory:
<https://computing.llnl.gov/tutorials/openMP/>
- Book: Rohit Chandra, et.al. „Parallel Programming in OpenMP"
Morgan Kaufmann, ISBN 1-55860-671-8



OpenMP - Overview:

C\$OMP FLUSH

#pragma omp critical

C\$OMP THREADPRIVATE (/ABC/)

CALL OMP_SET_NUM_THREADS(10)

C\$OMP parallel do shared(a, b, c)

call omp test lock(jlok)

- OpenMP: An API for Writing Multi-threaded Application
 - A set of compiler directives and library routines for parallel application programmers
 - Makes it easy to create multi-threaded (MT) programs in Fortran, C and C++
 - Standardizes last 15 years of SMP practice

#pragma omp parallel for private(A, B)

!\$OMP BARRIER

C\$OMP PARALLEL COPYIN (/blk/)

C\$OMP DO lastprivate(XX)

Nthrds = OMP_GET_NUM_PROCS()

omp_set_lock(lck)

OpenMP Overview: How is OpenMP typically used?

- OpenMP is usually used to parallelize loops:
 - Find your most time consuming loops.
 - Split them up between threads.

Split-up this loop between multiple threads

```
void main()
{
    double Res[1000];
    for( int i=0; i<1000; i++)
    {
        do_huge_comp(Res[i]);
    }
}
```

Sequential Program

```
#include "omp.h"
void main()
{
    double Res[1000];
    #pragma omp parallel for
    for(int i=0;i<1000;i++)
    {
        do_huge_comp(Res[i]);
    }
}
```

Parallel Program

Outline — Programming and Execution Model

- Introduction into OpenMP
- **Programming and Execution Model**
 - **Parallel regions: team of threads**
 - **Syntax**
 - **Data environment (part 1)**
 - **Environment variables**
 - **Runtime library routines**
 - **Exercise 1: Parallel region / library calls / privat & shared variables**
- Worksharing directives
 - Which thread executes which statement or operation?
 - Synchronization constructs, e.g., critical regions
 - Exercise 2: Pi
- Data environment and combined constructs
 - Nesting and Binding
 - Private and shared variables, Reduction clause
 - Combined parallel worksharing directives
 - Exercise 3: Pi with reduction clause and combined constructs
 - Exercise 4: Heat
- Summary of OpenMP API
- OpenMP Pitfalls & Optimization Problems

Components - Environment, Variables, Directives, Runtime

environment variables

directives

(special comment lines)

runtime library

```
#!/bin/tcsh
# Shell-Script
gcc -fopenmp test.c
setenv OMP_NUM_THREADS 4
a.out
```

```
/* Source file test.c */
#include <stdio.h>
#include <omp.h>
int main(void)
{
    #pragma omp parallel
    printf("me: %d\n", omp_get_thread_num());
}
```

```
me: 0
me: 3
me: 2
me: 1
```

OpenMP Overview: How do threads interact?

- OpenMP is a shared memory model.
 - Threads communicate by sharing variables.
- Unintended sharing of data causes race conditions:
 - **race condition**: when the program's outcome changes as the threads are scheduled differently.
- To control race conditions
 - Use synchronization to protect data conflicts.
- Synchronization is expensive so:
 - Change how data is accessed to minimize the need for synchronization.

OpenMP: Structured blocks

- Most OpenMP constructs apply to structured blocks
- Structured block: a block with one point of entry at the top and one point of exit at the bottom.
- The only “branches” allowed are `exit()` in C/C++.

```
#pragma omp parallel
{
    int id = omp_get_thread_num();
    more:
        res(id) = do_big_job(id);
        if( conv(res(id) ) goto more;
}
printf(" All done \n");
```

A structured block

```
if(go_now()) goto more;
#pragma omp parallel
{
    int id = omp_get_thread_num();
    more:
        res(id) = do_big_job(id);
        if( conv(res(id) ) goto done;
        go to more;
}
done:
    if(!really_done()) goto more;
```

Not a structured block

OpenMP: Structured blocks / Conditional Compilation

- In C/C++: a block is a single statement or a group of statements between brackets { }

```
#pragma omp parallel
{
    id = omp_thread_num();
    res(id) =
    lots_of_work(id);
}
```

```
#pragma omp for
for(I=0; I<N; I++)
{
    res[I] = big_calc(I);
    A[I] = B[I] + res[I];
}
```

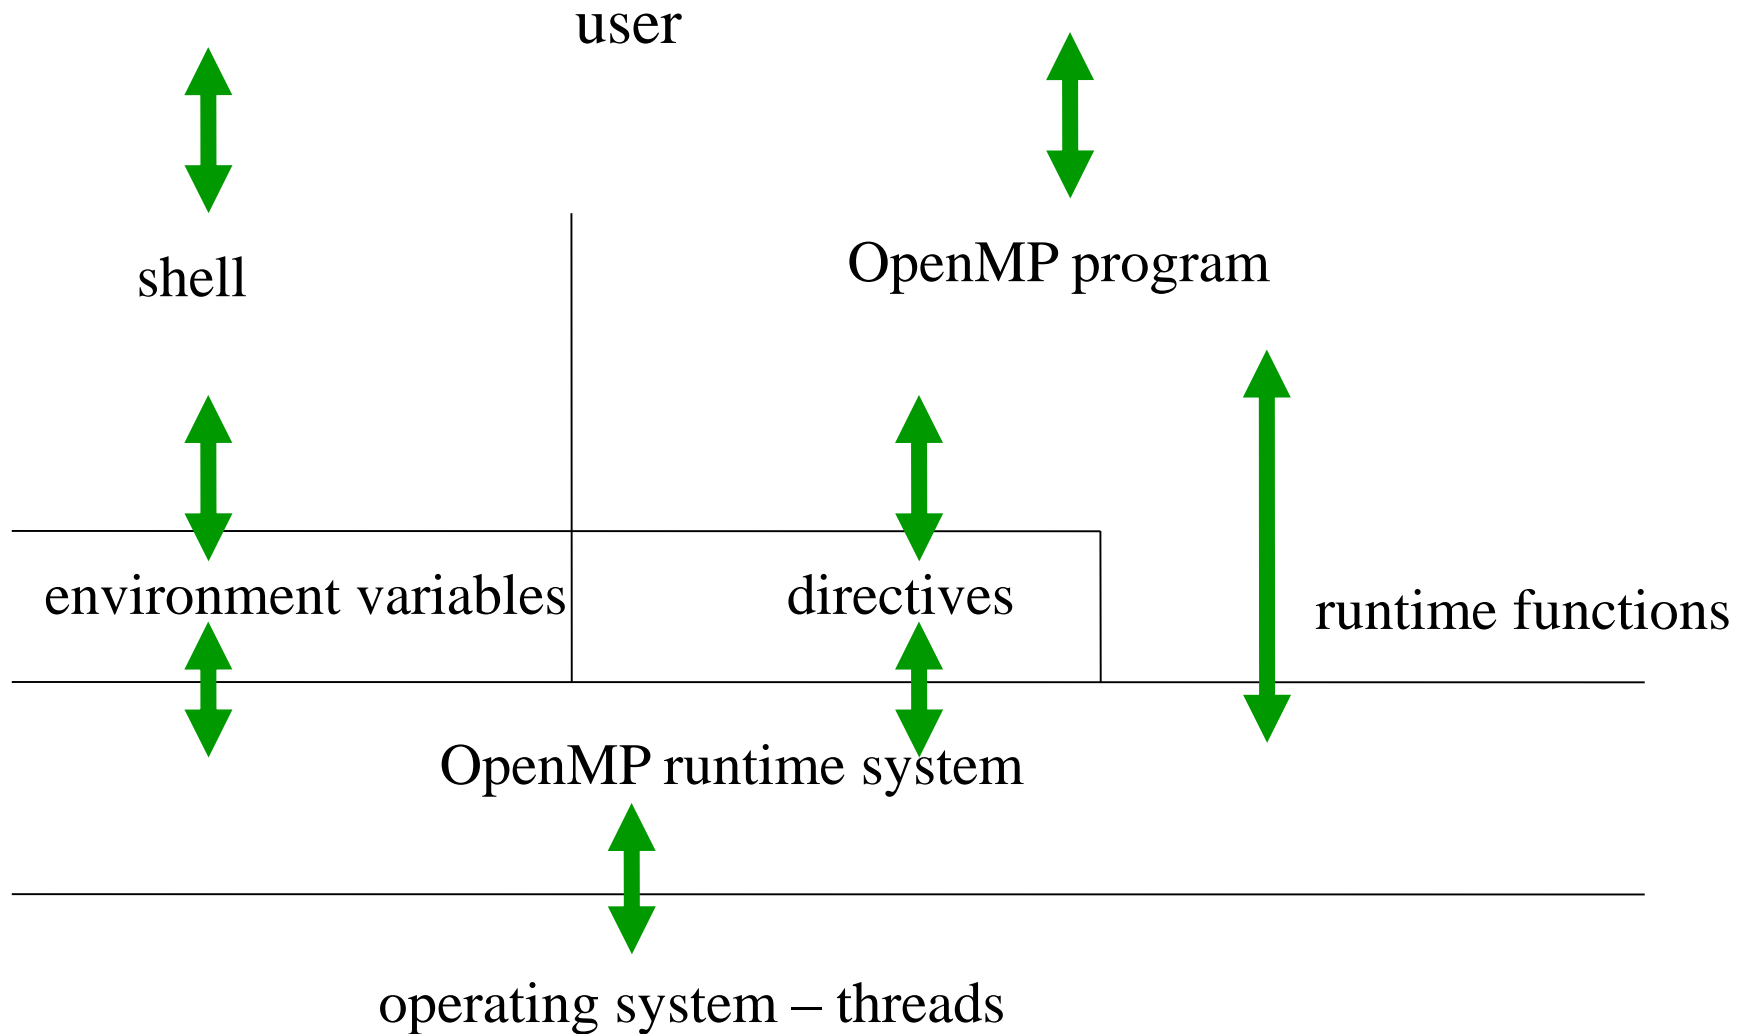
- Conditional Compilation

```
#ifdef _OPENMP
    iam=omp_get_thread_num();
#endif
```


```
/* OpenMP directive */
#pragma omp directive [clause ...]

/* OpenMP directive with
   continuation line */
#pragma omp directive clause \
    clause ...
```

OpenMP Components Diagram



Parallel Regions (1) - The fork-join concept

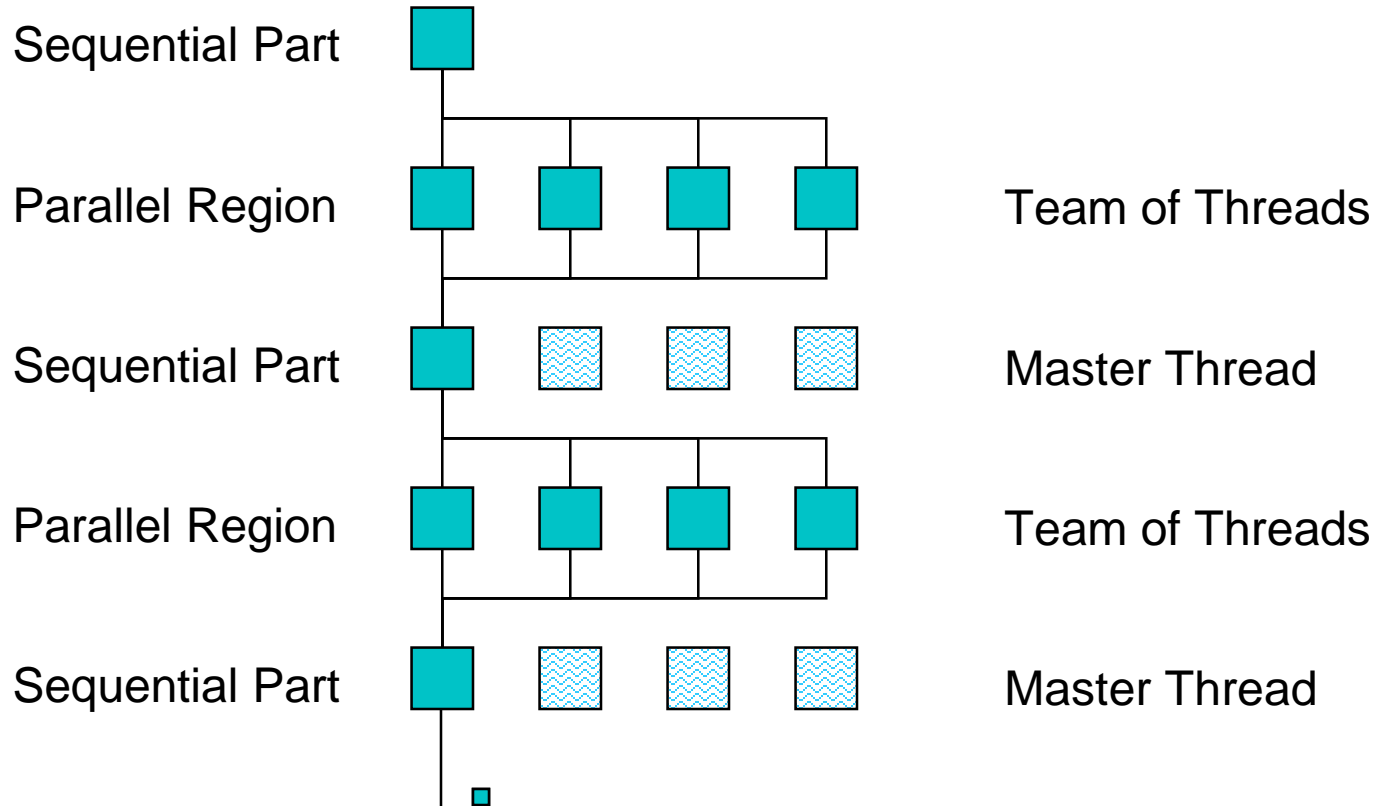


The OpenMP program starts like a serial program: single threaded

In the beginning of the first parallel region the slave threads are started. Together with the master, they form a team.

Between the parallel regions the slave threads are put to sleep.

OpenMP Execution Model




OpenMP: Some syntax details to get us started

- Most of the constructs in OpenMP are compiler directives or pragmas.
 - For C and C++, the pragmas take the form:
#pragma omp construct [clause [clause]...]
- Include OpenMP file
 - **#include "omp.h"**

C/C++

```
#pragma omp parallel  
    structured block  
/* omp end parallel */
```



OpenMP Parallel Region Construct Syntax

- C/C++:

```
#pragma omp parallel [ clause [ [ , ] clause ] ... ] new-line  
    structured-block
```

- *clause* can be one of the following:
 - private(*list*)
 - shared(*list*)
 - ...
- Conditional compilation

```
#ifdef _OPENMP  
    printf("%d avail.processors\n",omp_get_num_procs());  
#endif
```

- Include file for library routines:

```
#ifdef _OPENMP  
#include <omp.h>  
#endif
```

OpenMP Data Scope Clauses

- **private (list)**

Declares the variables in **list** to be private to each thread in a team

- **shared (list)**

Makes variables that appear in **list** shared among all the threads in a team

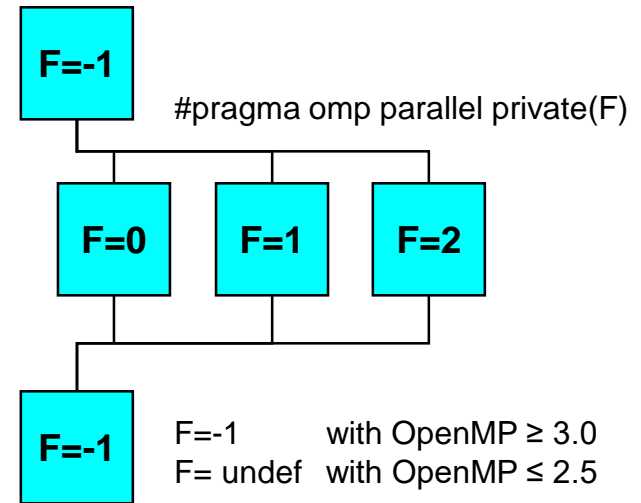
- If not specified: default shared, but

- stack (local) variables in called sub-programs are PRIVATE

- Loop control variable of parallel OMP

- for (C)

is PRIVATE



OpenMP Environment Variables

- **OMP_NUM_THREADS**

- sets the number of threads to use during execution
- when dynamic adjustment of the number of threads is enabled, the value of this environment variable is the maximum number of threads to use
- **setenv OMP_NUM_THREADS 16** [csh, tcsh]
- **export OMP_NUM_THREADS=16** [sh, ksh, bash]

- **OMP_SCHEDULE**

- applies only to **do/for** and **parallel do/for** directives that have the schedule type **RUNTIME**
- sets schedule type and chunk size for all such loops
- **setenv OMP_SCHEDULE "GUIDED,4"** [csh, tcsh]
- **export OMP_SCHEDULE="GUIDED,4"** [sh, ksh, bash]

Parallel Regions (2) - Runtime Library

```
#include "omp.h"
```

```
void main()
```

```
{
```

```
printf("inside parallel region? %d\n", omp_in_parallel());
```

```
printf("number of available processors? %d\n", omp_get_num_procs());
```

```
printf("maximum number of threads? %d\n", omp_get_max_threads());
```

```
omp_set_num_threads (omp_get_max_threads() );
```

```
#pragma omp parallel
```

```
{
```

```
printf("inside parallel region? %d\n", omp_in_parallel());
```

```
printf("number of threads in the team %d\n", omp_get_num_threads());
```

```
printf("my thread id %d\n", omp_get_thread_num() );
```

```
}
```

```
}
```

Sequent. region

Parallel region

```
mint [oschenk] export OMP_NUM_THREADS=3
```

```
mint [oschenk] ./a.out
```

```
inside parallel region? 0
```

```
number of available processors? 2
```

```
maximum number of threads? 3
```

```
inside parallel region? 1
```

```
number of threads in the team 3
```

```
my thread id 0
```

```
inside parallel region? 1
```

```
number of threads in the team 3
```

```
my thread id 2
```

```
inside parallel region? 1
```

```
number of threads in the team 3
```

```
my thread id 1
```

Parallel Regions (3) - Runtime Library

	Serial region	Parallel region
<code>void omp_set_num_threads (int)</code>	Set # threads to use in a team	don 't
<code>int omp_get_num_threads (void)</code>	1	Return # threads
<code>int omp_get_max_threads (void)</code>	Return max # threads (OMP_NUM_THREADS)	
<code>int omp_get_thread_num (void)</code>	0	Return thread id 0 ... #threads-1
<code>int omp_get_num_procs (void)</code>	Return # CPUs	
<code>void omp_set_dynamic (int)</code>	Control dynamic adjustment of # threads	don 't
<code>int omp_get_dynamic (void)</code>	.TRUE: if dynamic thread adjustment enabled .FALSE. Otherwise	
<code>int omp_in_parallel (void)</code>	.FALSE.	.TRUE.

Parallel Regions (4)

- You create threads in OpenMP with the “omp parallel” pragma.
- For example, to create a 4 thread parallel region:

Each thread executes a copy of the code within the structured block

```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
  int ID = omp_get_thread_num();  
  pooh(ID,A);  
}
```

Runtime function to request a certain number of threads

Runtime function returning a thread ID

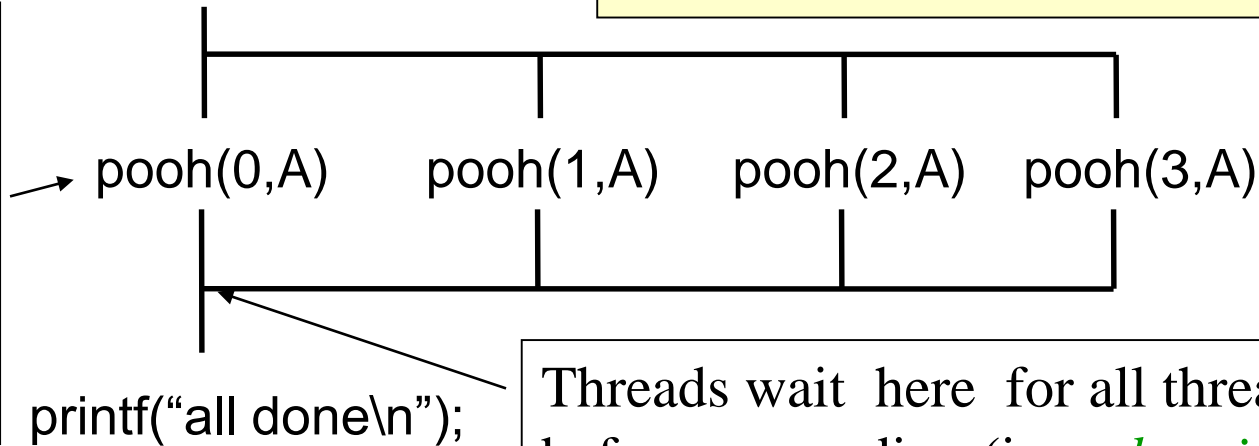
- Each thread calls pooh(ID,A) for ID = 0 to 3

Parallel Regions (5)

- Each thread executes the same code redundantly.

```
double A[1000];  
  
omp_set_num_threads(4)
```

A single copy of A is shared between all threads.



```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
    int ID = omp_get_thread_num();  
    pooh(ID, A);  
}  
printf("all done\n");
```

Threads wait here for all threads to finish before proceeding (i.e. a *barrier*)

OpenMP Runtime Library (3): Wall clock timers OpenMP 2.0

- Portable wall clock timers
- **DOUBLE PRECISION FUNCTION OMP_GET_WTIME()**
provides elapsed time

```
START= OMP_GET_WTIME()  
! Work to be measured  
END = OMP_GET_WTIME()  
printf("Work took %e seconds\n", END-START);
```

- **DOUBLE PRECISION FUNCTION OMP_GET_WTICK()**
returns the number of seconds between two successive clock ticks

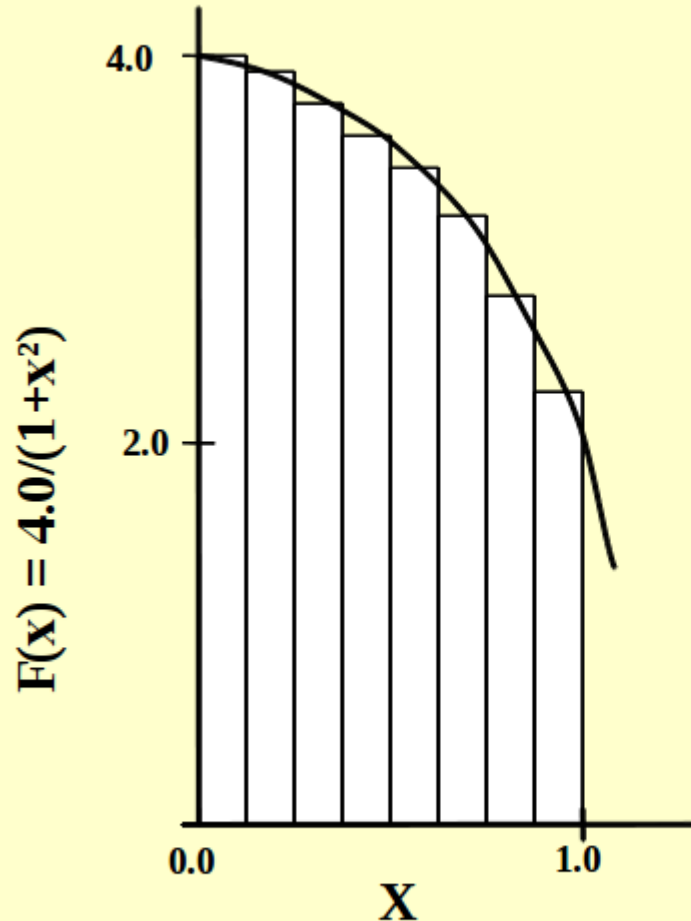
Outline

- Introduction into OpenMP
- Programming and Execution Model
 - Parallel regions: team of threads
 - Syntax
 - Data environment (part 1)
 - Environment variables
 - Runtime library routines
 - **Exercise 1: Parallel region / library calls / privat & shared variables**
- Worksharing directives
 - Which thread executes which statement or operation?
 - Synchronization constructs, e.g., critical regions
 - Nesting and Binding
 - Exercise 2: Pi
- Data environment and combined constructs
 - Private and shared variables, Reduction clause
 - Combined parallel worksharing directives
 - Exercise 3: Pi with reduction clause and combined constructs
 - Exercise 4: Heat
- Summary of OpenMP API
- OpenMP Pitfalls & Optimization Problems

In-class exercise: A multi-threaded “pi” program

- On the following slide, you’ll see a sequential program that uses numerical integration to compute an estimate of PI.
- Parallelize this program using OpenMP.
- Remember, you’ll need to make sure multiple threads don’t overwrite each other’s private variables.

In-class exercise: A multi-threaded “pi” program



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .

In-class exercise: Pi program: The sequential program

```
static long num_steps = 100000;
double step;
void main ()
{
    int i;
    double x, pi, sum = 0.0;

    step = 1.0/(double) num_steps;

    for (i=0; i<num_steps; i++)
    {
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x) ;
    }
    pi = step * sum;
}
```

In-class exercise: pi0.c – sequential code

The include and timing blocks are removed on the next slides

```
#include <stdio.h>
#include <time.h>
#include <sys/time.h>
#ifdef _OPENMP
# include <omp.h>
#endif
include block
#define f(A) (4.0/(1.0+A*A))
const int n = 10000000;
```

```
int main(int argc, char** argv)
{
    int i;
    double w,x,sum,pi;
```

```
    clock_t t1,t2;
    struct timeval tv1,tv2;
    struct timezone tz;
# ifdef _OPENMP
    double wt1,wt2;
# endif
timing block A
```

```
    gettimeofday(&tv1, &tz);
# ifdef _OPENMP
    wt1=omp_get_wtime();
# endif
    t1=clock(); timing block B
```

```
/* pi = integral [0..1] 4/(1+x**2) dx */
w=1.0/n;
sum=0.0;
for (i=1;i<=n;i++)
{
    x=w*((double)i-0.5);
    sum=sum+f(x);
}
pi=w*sum; the calculation
```

```
    t2=clock(); timing block C
# ifdef _OPENMP
    wt2=omp_get_wtime();
# endif
    gettimeofday(&tv2, &tz);
    printf( "computed pi = %24.16g\n", pi);
    printf( "CPU time (clock)
= %12.4g sec\n", (t2-t1)/1000000.0 );
# ifdef _OPENMP
    printf( "wall clock time
(omp_get_wtime) = %12.4g sec\n",
    wt2-wt1 );
# endif
    printf( "wall clock time (gettimeofday)
= %12.4g sec\n",
    (tv2.tv_sec-tv1.tv_sec) +
    (tv2.tv_usec-tv1.tv_usec)*1e-6 );
    return 0;
}
```

In-class exercise: Parallel region (1)

- Goal: usage of
 - runtime library calls
 - conditional compilation, environment variables
 - parallel regions, private and shared clauses
- Serial programs: **pi0.c** on <https://www2.icorsi.ch/>
- Build a team of 2 students, compile **serial** program **pi0.c** on CUB cluster and run e.g. with
 - **scp pi0.c student@cub.inf.usi.ch:.**
 - **ssh student@cub.inf.usi.ch**
 - **salloc -pdebug -t 00:30:00** (to request one node in interactive mode for 30 min)
- Compile **as OpenMP** program and run on 4 core
 - **gcc -O3 -fopenmp pi0.c -o pi**
 - **export OMP_NUM_THREADS=4**
 - **./pi**

expected result: program is not parallelized,
 therefore same pi-value and timing,
 additional output from omp_get_wtime()

In-class exercise: pi1.c

- Modify pi0.c -> pi1.c
- Directly after the declaration part, add in a **parallel region that prints on each thread**
 - **its rank** (with **omp_get_thread_num()**) and
 - **the number of threads** (with **omp_get_num_threads()**)
- compile **gcc -O3 -fopenmp pi1.c -o pi1** and run on 4 cores
- Expected results: numerical calculation is still not parallelized, therefore still same pi-value and timing, additionally output:

```
bash$ gcc -O3 -fopenmp -o pi1 pi1.c
bash$ export OMP_NUM_THREADS=4; ./pi1
I am thread 0 of 4 threads
I am thread 2 of 4 threads
I am thread 3 of 4 threads
I am thread 1 of 4 threads
computed pi = 3.141592653589731
CPU time (clock) = 0.16 sec
wall clock time (omp_get_wtime) = 0.1681 sec
wall clock time (gettimeofday) = 0.1681 sec
```

} undefined sequence!

OpenMP Advanced Exercise: pi1.c

- Modify pi1.c
- Use a private variable for the rank of the threads
- Check, whether you can get a race-condition if you forget the private clause on the **omp parallel** directive, e.g.

I am thread	2	of	4	threads
I am thread	2	of	4	threads
I am thread	2	of	4	threads
I am thread	2	of	4	threads

- Don't wonder if you get always correct output because the compiler may use on each thread a private register instead of writing into the shared memory

OpenMP Advanced Exercise: pi1.c

- Modify pi1.c
- Guarantee with conditional compilation, that source code still works **with non-OpenMP compilers** (i.e., without OpenMP compile-option).
- Add an “**else clause**”, printing a text if OpenMP is not used.
- Expected output:

If compiled with OpenMP, see previous slide.

If compiled without OpenMP

```
bash$ gcc -o pi1 pi1.c
bash$ export OMP_NUM_THREADS=4; ./pi1
This program is not compiled with OpenMP
computed pi =          3.1415926535897931
CPU time (clock)          =          0.16 sec
wall clock time (gettimeofday) =          0.1681 sec
```

Outline

- Introduction into OpenMP
- Programming and Execution Model
 - Parallel regions: team of threads
 - Syntax
 - Data environment (part 1)
 - Environment variables
 - Runtime library routines
 - Exercise 1: Parallel region / library calls / privat & shared variables
- **Worksharing directives**
 - **Which thread executes which statement or operation?**
 - **Synchronization constructs, e.g., critical regions**
 - **Nesting and Binding**
 - **Exercise 2: Pi**
- Data environment and combined constructs
 - Private and shared variables, Reduction clause
 - Combined parallel worksharing directives
 - Exercise 3: Pi with reduction clause and combined constructs
 - Exercise 4: Heat
- Summary of OpenMP API
- OpenMP Pitfalls & Optimization Problems

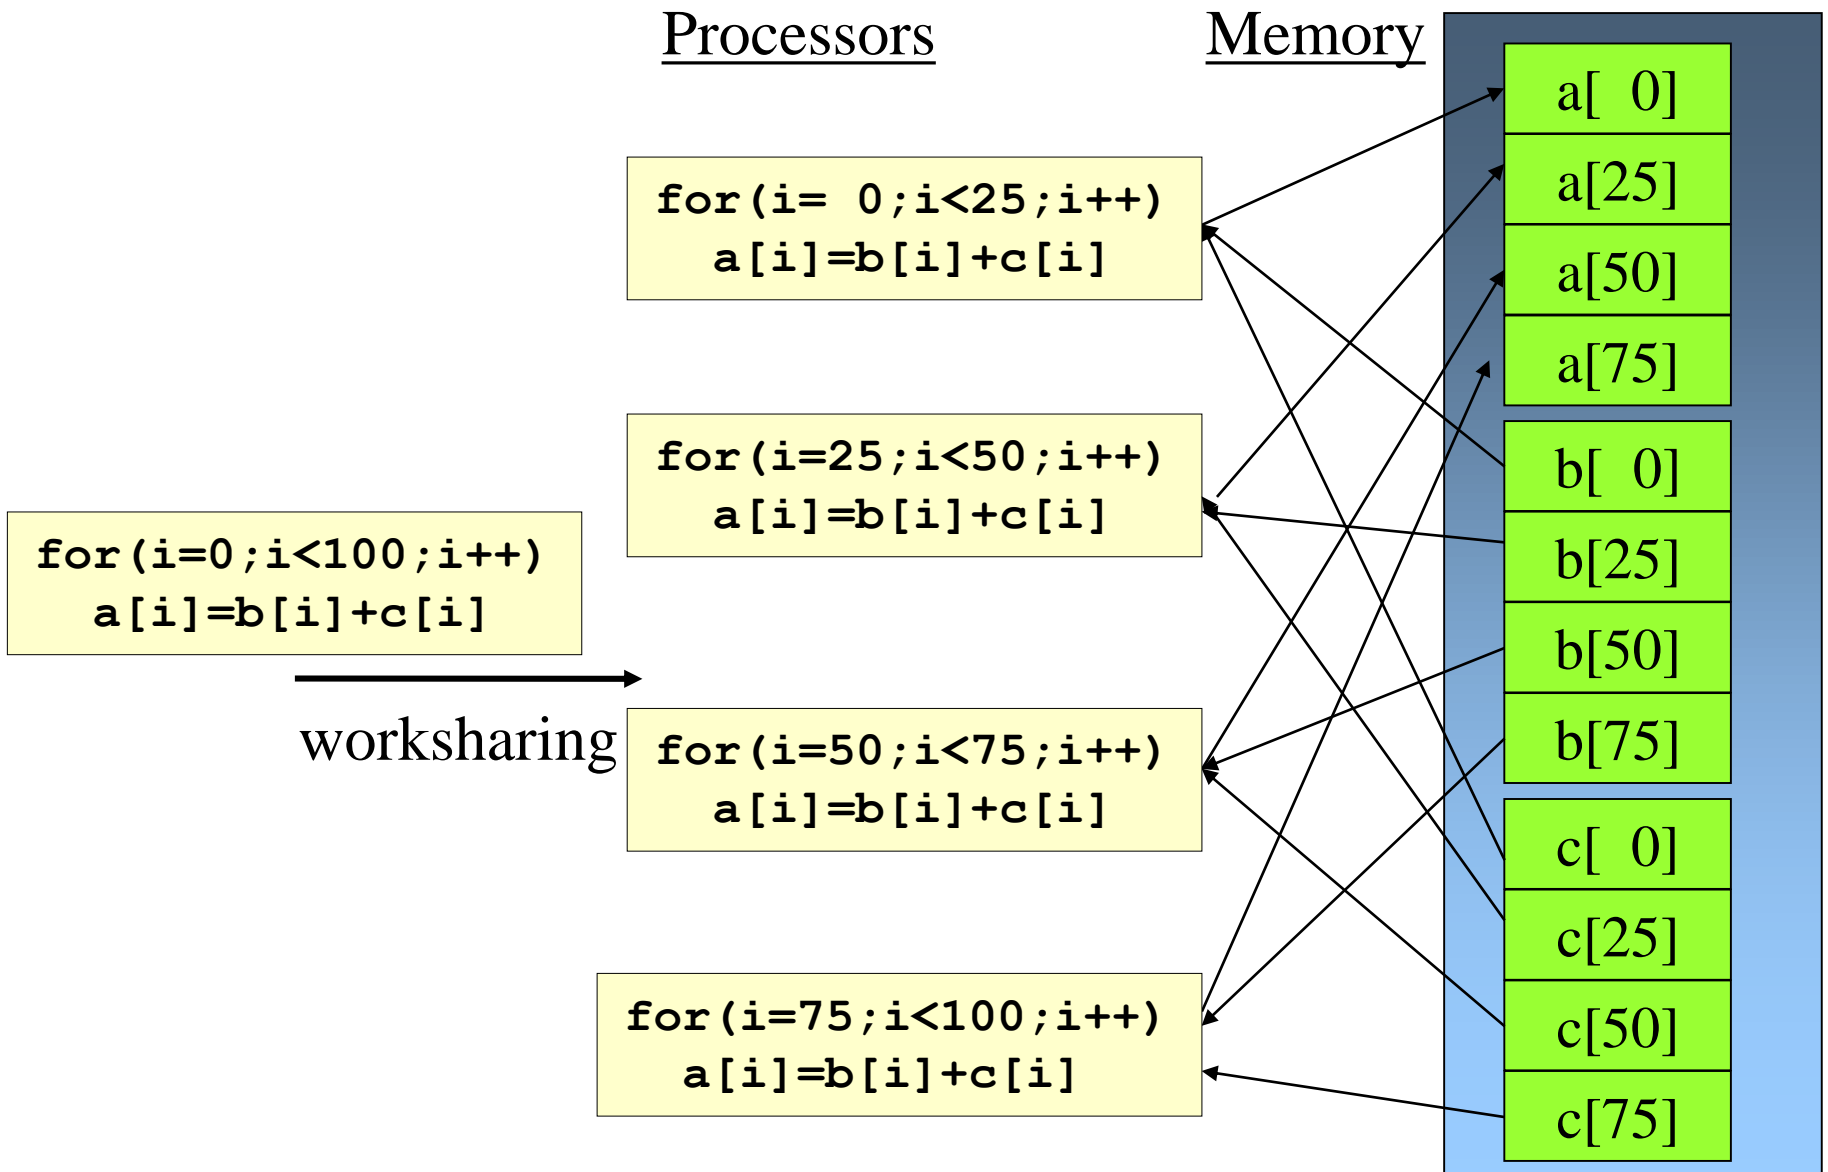
Worksharing and Synchronization

- Which thread executes which statement or operation?
- and when?
 - Worksharing constructs
 - Master and synchronization constructs
- **i.e., organization of the parallel work!!!**

Worksharing and Synchronization

- Divide the execution of the enclosed code region among the members of the team
- Must be enclosed dynamically within a parallel region
- They do not launch new threads
- No implied barrier on entry
 - **for** directive
 - **sections** directive
 - **task** directive
 - **single** directive

Work Sharing (1) - Principle



Work Sharing (1) - A motivating example

Sequential code

```
for(i=0;I<N;i++)    { a[i] = a[i] + b[i];}
```

OpenMP
parallel region

```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    for(i=istart;I<iend;i++)
    { a[i] = a[i] + b[i];}
}
```

OpenMP
parallel region and
a work-sharing
for-construct

```
#pragma omp parallel
#pragma omp for schedule(static)
for(i=0;I<N;i++)    { a[i] = a[i] + b[i];}
```

Work Sharing (2) - Sharing Constructs

- The “for” **work-sharing** construct splits up loop iterations among the threads in a team

```
#pragma omp parallel  
  
for (I=0; I<N; I++)  
{  
    NEAD_STUFF(I);  
}
```

All threads select all
loop indices “I”

```
#pragma omp parallel  
#pragma omp for  
for (I=0; I<N; I++)  
{  
    NEAD_STUFF(I);  
}
```

The loop indices “I”
are distributed
among threads

By default, there is a **barrier** at the end of the “omp for”. Use the **nowait** clause to turn off the barrier.

Work Sharing (7) - Combined parallel/work-share

- OpenMP shortcut:
 - Put the “parallel” and the work-share on the same line

```
double  res[MAX];  
int i;  
#pragma omp parallel  
{  
    #pragma omp for  
    for (i=0;i< MAX; i++)  
    {  
        res[i] = huge();  
    }  
}
```

```
double  res[MAX];  
int i;  
#pragma omp parallel for  
for (i=0;i< MAX; i++)  
{  
    res[i] = huge();  
}
```


These are equivalent

The diagram consists of two code snippets, one on the left and one on the right, each enclosed in a yellow box. Below these boxes is a white rectangular box containing the text "These are equivalent". Two black arrows originate from this box: one points to the line "#pragma omp for" in the left snippet, and the other points to the line "#pragma omp parallel for" in the right snippet.

OpenMP **for** Directives – Syntax

- Immediately following loop executed in parallel
#pragma omp for [*clause*[[,]*clause*] ...] *new-line*
for-loop
- The corresponding **for** loop must have *canonical shape*:

```
for( [integer type] var=lb; var < b; var++      )  
    <=      ++var  
    >      var+=incr  
    >=     var=var+incr  
    var-- ...
```

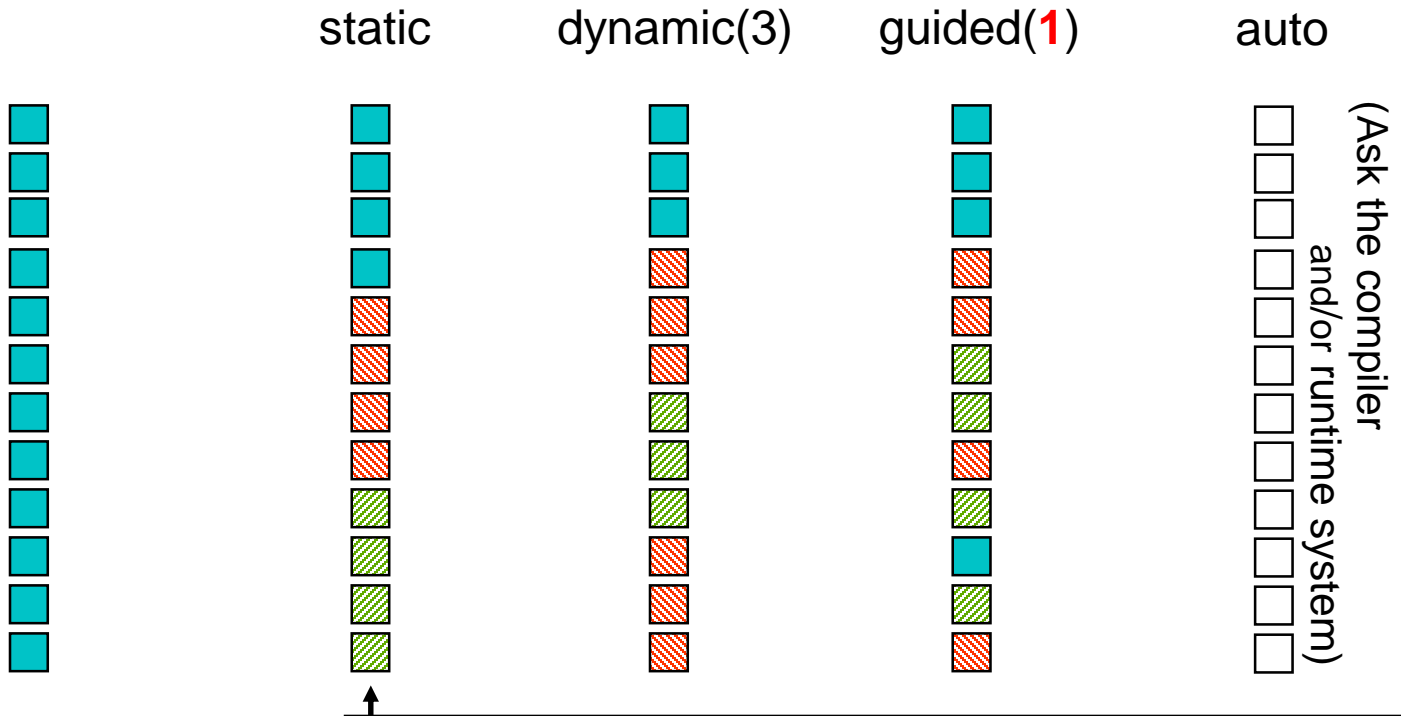
`var`: must not be modified in the loop body;
integer (signed or **unsigned**),
or **pointer type** (C only),  (**OpenMP ≥ 3.0**)
or **random access iterator type** (C++ only)

`lb, b, incr`: loop invariant expression
→ the number of iterations must be computable at loop begin

Work Sharing (3) - The schedule clause

- The schedule clause effects how loop iterations are mapped onto threads
 - **schedule(static [,chunk])**
 - Deal-out blocks of iterations of size “chunk” to each thread.
 - In our example: thread #0: i=0 to 24; thread #1: i=25 to 49; ...
 - **schedule(dynamic[,chunk])**
 - Each thread grabs “chunk” iterations off a queue until all iterations have been handled.
 - In our example: thread #0: i=0, 3, 8, ..; thread #1: i=1, 2, 5, ..;
 - **schedule(guided[,chunk])**
 - Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds.
 - **schedule(auto)**
 - Scheduling is delegated to the compiler and/or runtime system
(OpenMP ≥ 3.0)
 - **schedule(runtime)**
 - Schedule and chunk size taken from the OMP_SCHEDULE environment variable.


Loop scheduling



- Method is implementation dependent, e.g.,
13 iterations on 3 threads = $5+5+3$ or $= 5+4+4$
 - Two loops in same parallel region and
with same count are divided in same way,
i.e., static schedule is **deterministic**
- } **OpenMP ≥ 3.0**

Work Sharing (4) – Matrix Transposition

```
#pragma omp parallel for private(h,i,j)
schedule static numthreads(3)
for (i=0; i<n; i++) {
    for (j=i+1; j<n; j++) {
        h = A[i][j];
        A[i][j] = A[j][i];
        A[j][i] = h;
    }
} // end parallel for
```



- Thread 0 would have much more work than thread 2!

Work Sharing (5) - The schedule clause

Schedule Clause	When To Use
STATIC	Predictable and similar work per iteration thread #0: i=0 to24; thread #1: i=25 to 49 thread #2: i=50 to74; thread #0: i=75 to 99
DYNAMIC	Unpredictable, highly variable work per iteration thread #0: i=0,2, 8,..., thread #1: 4,6,7,... thread #2: i=2,3,9,..; thread #0: i=15,18,91,..
GUIDED	Special case of dynamic to reduce scheduling overhead

Work Sharing (6) - The section clause

- The **section work-sharing** construct gives a different structured block to each thread.

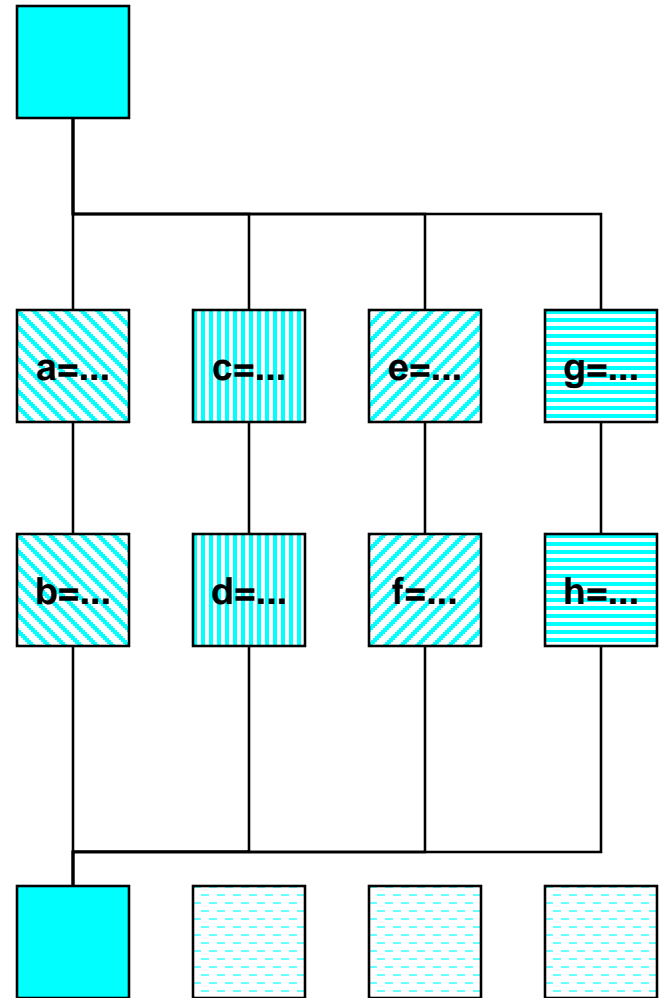
```
#pragma omp parallel
#pragma omp sections
{
    #pragma omp section
    x_calculation(); // only one thread
    #pragma omp section
    y_calculation(); // only one thread
    #pragma omp section
    z_calculation(); // only one thread
}
```

By default, there is a **barrier** at the end of the “omp sections”.
Use the “**nowait**” clause to turn off the barrier.

Work Sharing (6) - The **section** clause

C / C++:

```
#pragma omp parallel
{
  #pragma omp sections
  {
    #pragma omp section
    {
      a=...;
      b=...;
    }
    #pragma omp section
    {
      c=...;
      d=...;
    }
    #pragma omp section
    {
      e=...;
      f=...;
    }
    #pragma omp section
    {
      g=...;
      h=...;
    }
  } /*omp end sections*/
} /*omp end parallel*/
```



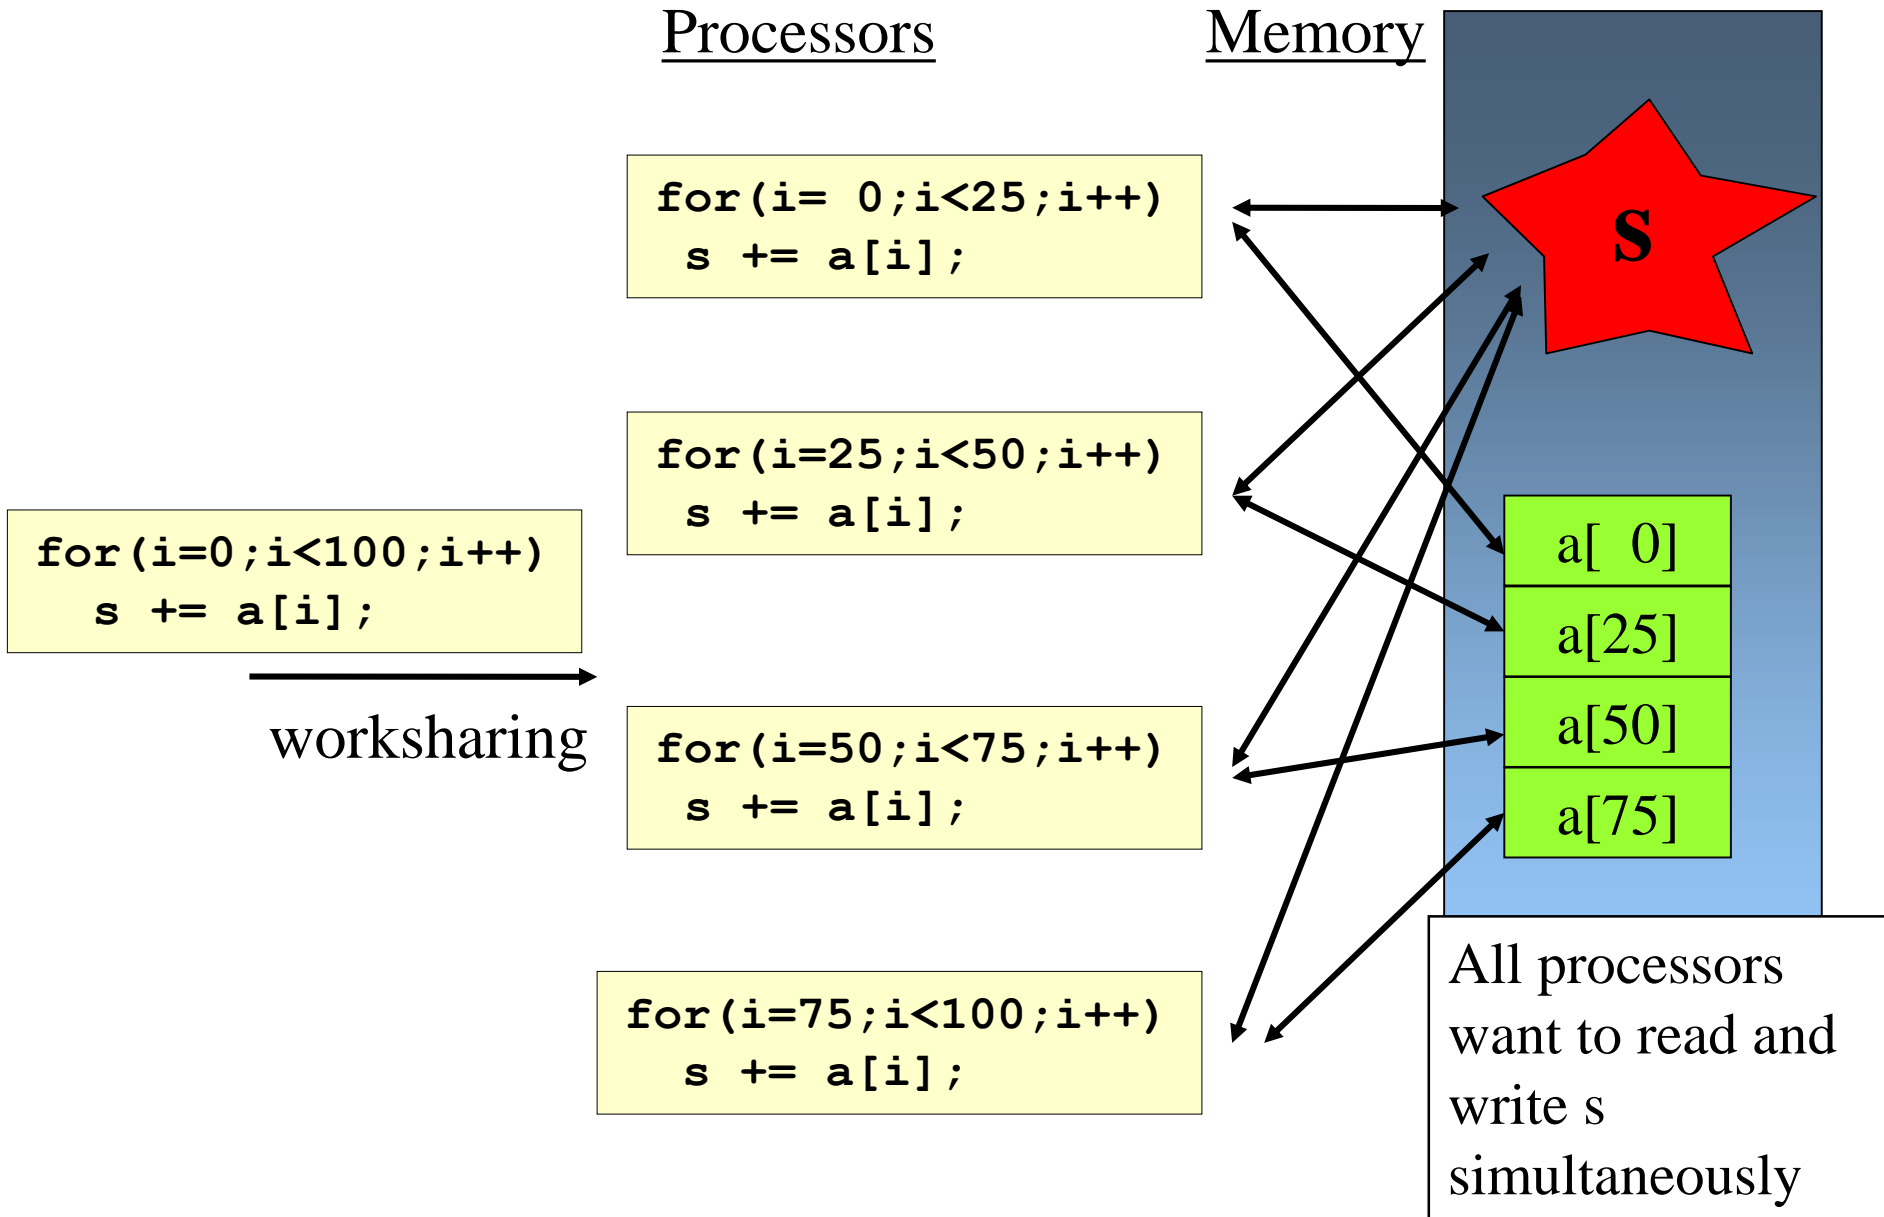
OpenMP **task** Directive – Parallelized traversing of a tree

C/C++

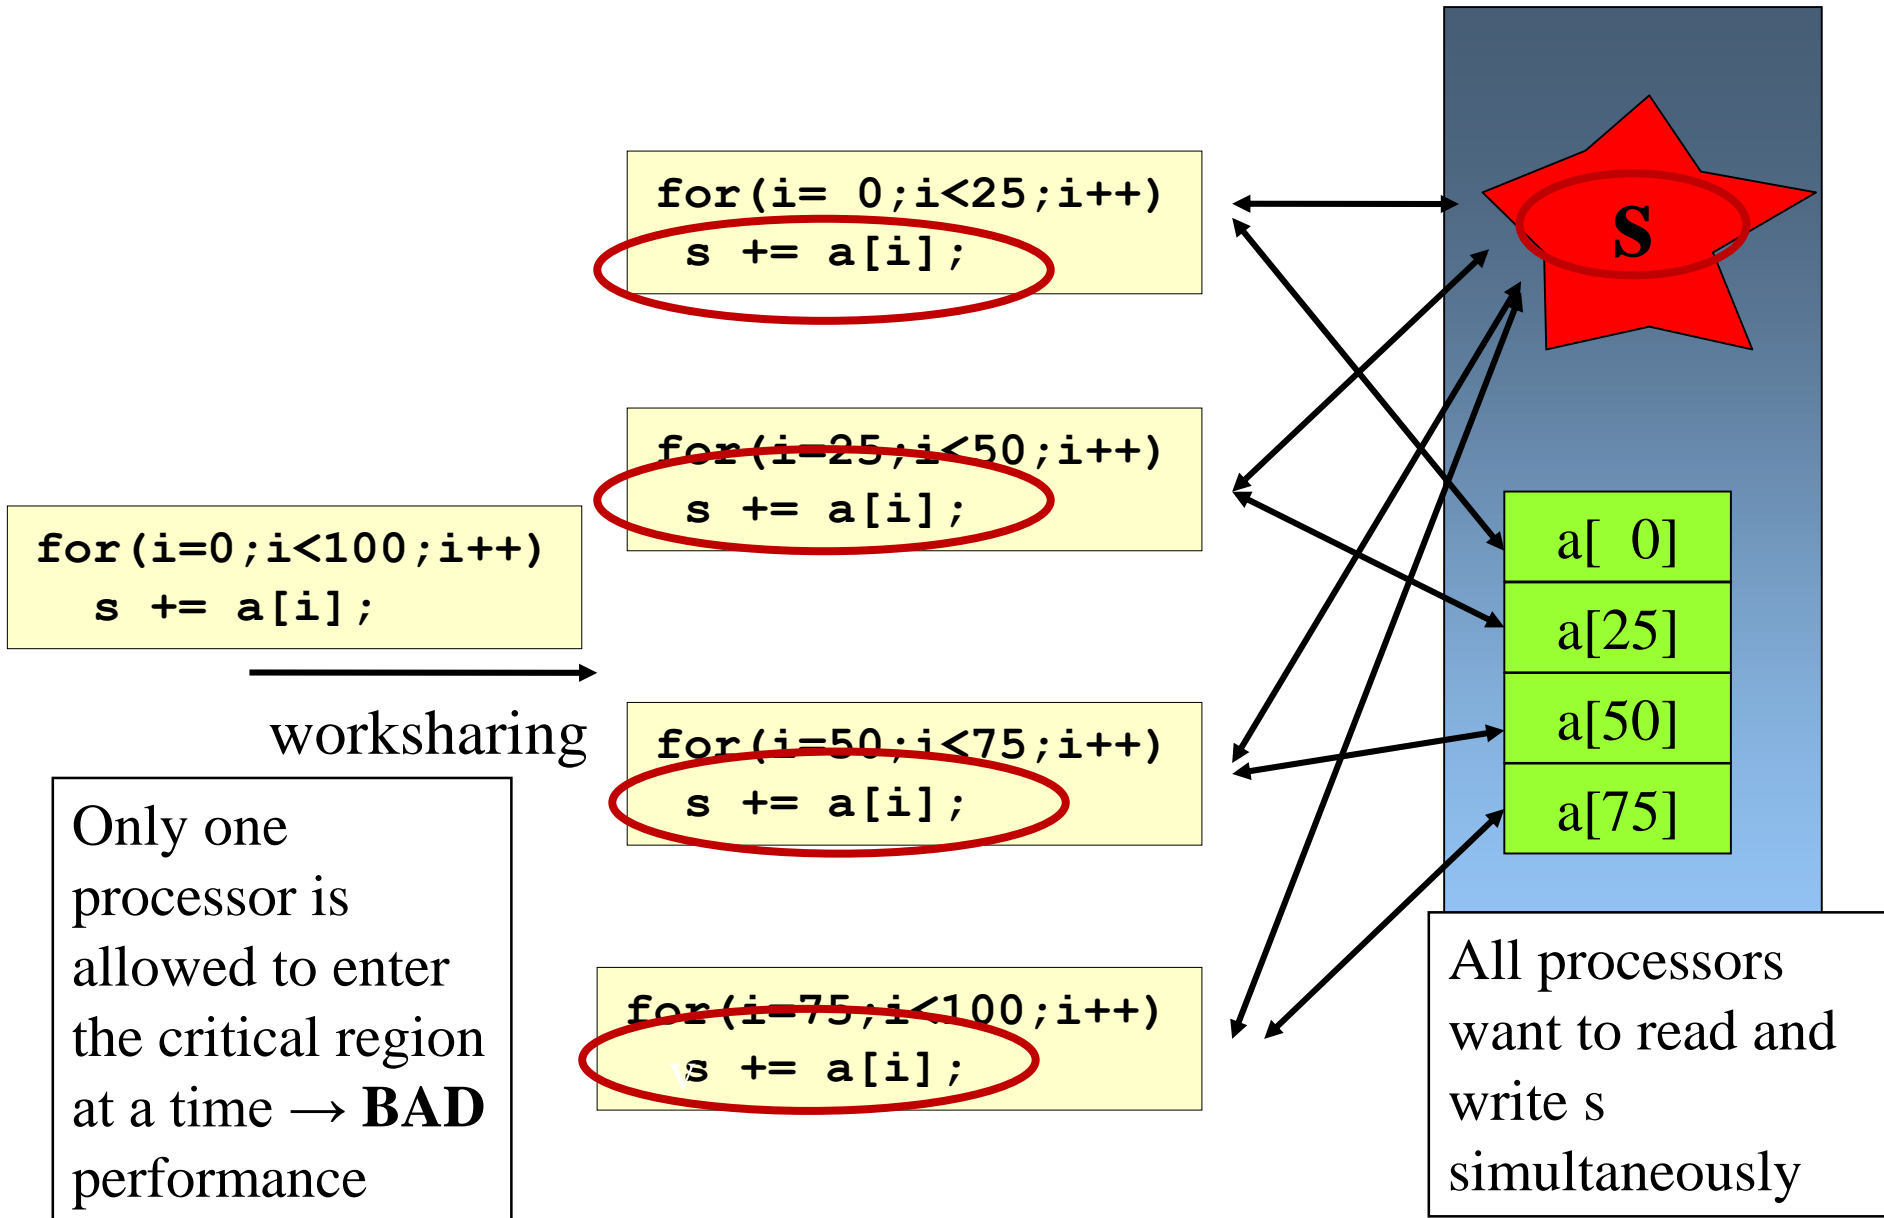
```
struct node {
    struct node *left;
    struct node *right;
};
extern void process(struct node *);
void traverse( struct node *p ) {
    if (p->left)
#pragma omp task // p is firstprivate by default
        traverse(p->left);
    if (p->right)
#pragma omp task // p is firstprivate by default
        traverse(p->right);
    process(p); // significant work with p
}
int main(int argc, char **argv)
{ struct node tree;
  ... // producing the tree
#pragma omp parallel
  {
#pragma omp single
  {
      traverse(&tree); //traversing the existing tree
  } // end of omp single
  } // end of omp parallel
}
```

- Starting the parallel team of threads
- Using only one thread for starting the traversal
- First execution with single thread (= 1st task)
- A new task is started (on a new thread)
- A recursive call to traverse() in this 2nd task
- 3rd task is started
- Work is done in 1st task
- Recursive calls starting 4th, 5th, ... tasks

Critical sections (1)




Critical sections (2)



Critical sections (3)

- Only one thread at a time can enter a **critical** section.

Threads wait their turn – only one at a time calls `consum()`



```
float res;  
#pragma omp parallel  
{  
    float B;  
    int i;  
    #pragma omp for  
    for(i=0;i<niters;i++)  
    {  
        B = big_job(i);  
        #pragma omp critical  
        consum (B, RES);  
    }  
}
```

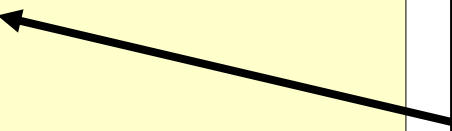

Critical sections (4) – Critical / end critical

```
for(i=0;i<100;i++)  
    s = s + a[i];
```

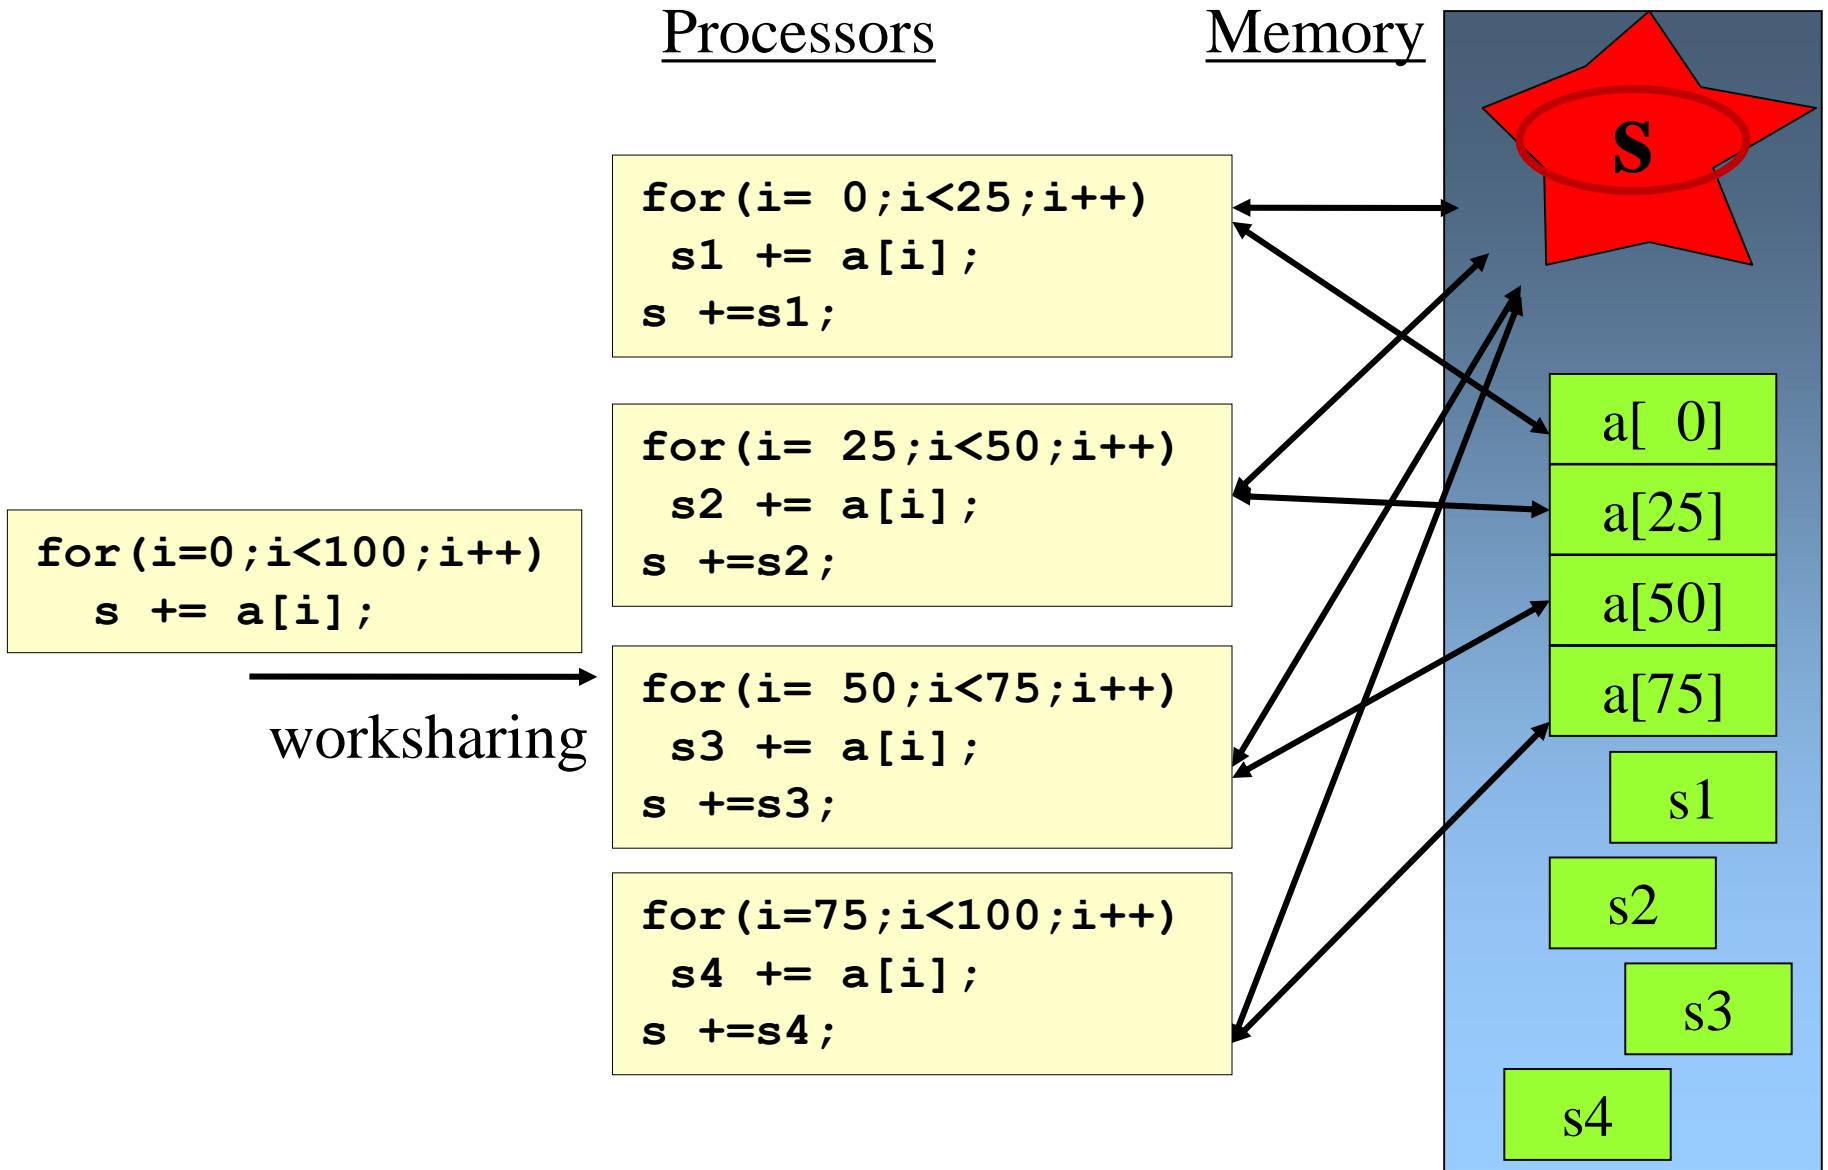
```
#pragma omp parallel for private (i)  
for (i=0; i<100; i++) {  
    #pragma omp critical  
    { s += a[i]; }  
}
```

Only one processor is allowed to enter the **critical section** at a time

As the loop body consists of a critical section only, the parallel program will run much slower



Critical section (5)



Critical section (6) – Critical / end critical

```
for(i=0;i<100;i++)  
    s = s + a[i];
```

```
#pragma omp parallel for \  
    private (i)  
for (i=0; i<100; i++) {  
    #pragma omp critical  
    { s += a[i]; }  
}
```

```
#pragma omp parallel \  
    private (i,s_local)  
{  
    s_local = 0;  
    #pragma omp for  
    for (i=0; i<100; i++)  
        { s_local += a[i]; }  
    #pragma omp critical  
        { s += s_local; }  
}
```

Only one processor is allowed to enter the **critical section** at a time

As the loop body consists of a critical region only, the parallel program will run much slower

Now the partial sums are calculated in parallel. The critical region is entered only once per thread.

Outline

- Introduction into OpenMP
- Programming and Execution Model
 - Parallel regions: team of threads
 - Syntax
 - Data environment (part 1)
 - Environment variables
 - Runtime library routines
 - Exercise 1: Parallel region / library calls / privat & shared variables
- Worksharing directives
 - Which thread executes which statement or operation?
 - Synchronization constructs, e.g., critical section
 - **Exercise 2: Pi**
- Data environment and combined constructs
 - Nesting and Binding
 - Private and shared variables, Reduction clause
 - Combined parallel worksharing directives
 - Exercise 3: Pi with reduction clause and combined constructs
 - Exercise 4: Heat
- Summary of OpenMP API
- OpenMP Pitfalls & Optimization Problems

In-class exercise 2: pi Program (1)

- Goal: usage of
 - worksharing constructs: `#pragma omp for`
 - `#pragma omp critical` directive
 - Use your result `pi1.c` from the last in-class exercise
 - Modify `pi1.c` -> `pi2.c`
 - compile serial program `pi2.c` and run
- add `parallel region` and `pragma omp for` directive and compile
- set environment variable `OMP_NUM_THREADS` to `2` and run
value of pi? (should be wrong!)
- run again
value of pi? (...wrong and **unpredictable**)
- set environment variable `OMP_NUM_THREADS` to `4` and run
value of pi? (...and stays wrong)
- run again
value of pi? (...but where is the race-condition?)

In-class exercise 2: pi Program (2)

- add **private (x)** clause in **pi2.c** and compile
- set environment variable **OMP_NUM_THREADS** to **2** and run
value of pi? (should be still incorrect ...)
- run again
value of pi?
- set environment variable **OMP_NUM_THREADS** to **4** and run
value of pi?
- run again
value of pi? (... and where is the second race-condition?)

In-class exercise 2: pi Program (3)

- add **critical** directive in **pi2.c** around the sum-statement and compile
- set environment variable **OMP_NUM_THREADS** to **2** and run
value of pi? (should be now correct!, but huge CPU time!)
- run again
value of pi? (but not reproducible in the last bit!)
- set environment variable **OMP_NUM_THREADS** to **4** and run
value of pi? execution time? (Oh, does it take longer?)
- run again
value of pi? execution time?
How can you optimize your code?

In-class advanced exercise 2: pi Program (4)

- Modify the printing of the thread rank and the number of threads from Exercise 1:
 - Only one thread should print the real number of threads used in parallel regions.
 - For this, use a **single** construct
 - Expected result:

OpenMP-parallel with 4 threads

```
computed pi =          3.14159265358967
CPU time (clock)          =          0.01659 sec
wall clock time (omp_get_wtime) =          0.01678 sec
wall clock time (gettimeofday) =          0.01679 sec
```